

# Operating Systems

## Tutorial

Fabian Klopfer

November 2, 2020

# TOC

Manuals & Literature

Operating Systems Overview

Systems Programming Overview

Compilers & Linking

References

# Intro

- ▶ Fabian, 2. Semester M.Sc. Computer & Information Science,  
fabian.klopfen@uni.kn
- ▶ Reminder: English Poll
- ▶ Fork [the repository](#) & do sheet 0
- ▶ More organisation stuff in the lecture
- ▶ The following slides are preliminary as the lecture is WIP
- ▶ sorry for the sloppy image references!

# Processor & Compiler Reference

Not a must read for this course but a good reference for low level details:

[AMD64 Manual](#)

[Intel x86 & x64 Manual](#)

[SystemV x86 Calling Convention](#)

[Comprehensive Low Level Information Collection](#)

# C Programming & Operating System Reference

Your best friend when programming:

[Official C Language Reference](#)

Unix & Linux references:

[Linux Manuals \(searchable!\)](#)

[Linux Syscalls Man Page](#)

[POSIX Specification](#)

[The Linux Kernel Documentation Project](#)

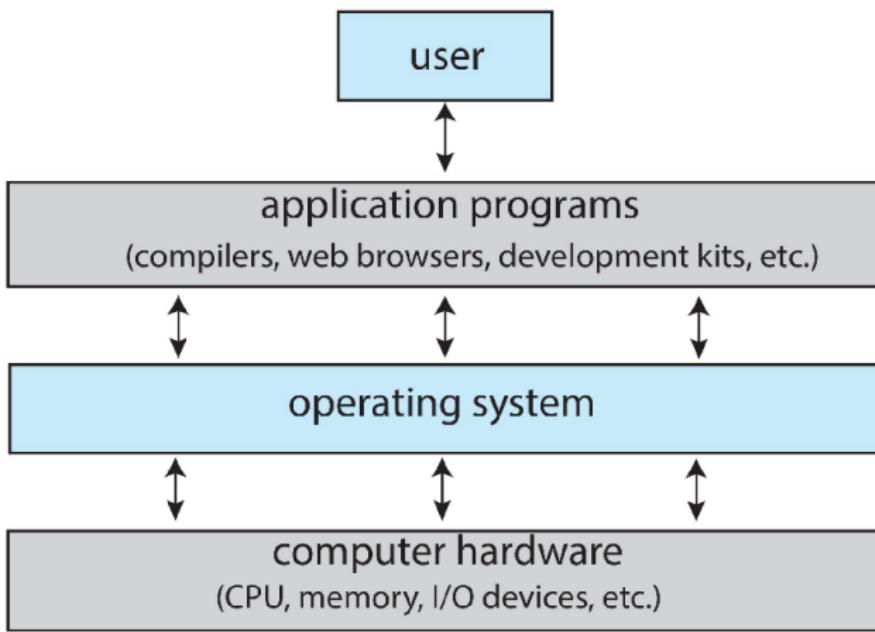
## Literature hint

Get (digital) copies of all books!

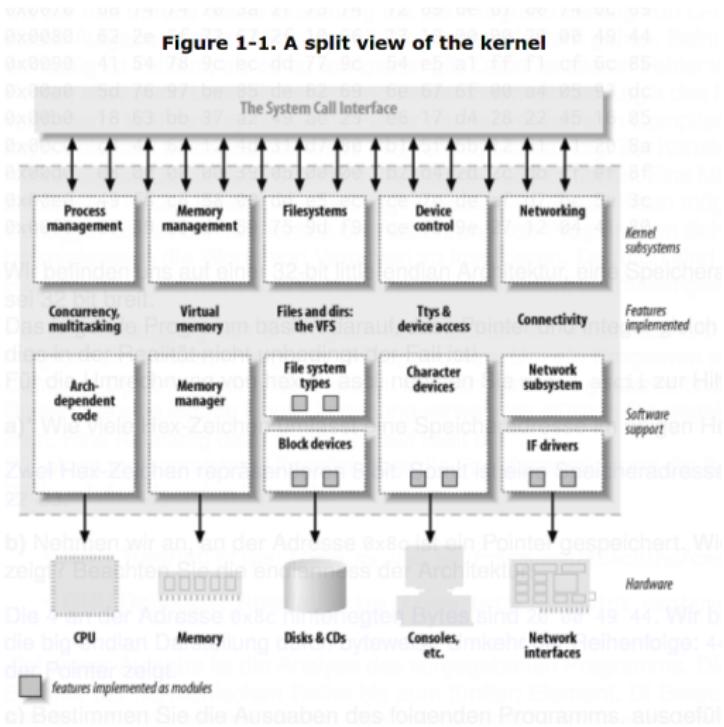
Many nice graphs:

[https://dinus.ac.id/repository/docs/ajar/Operating\\_System.pdf](https://dinus.ac.id/repository/docs/ajar/Operating_System.pdf)

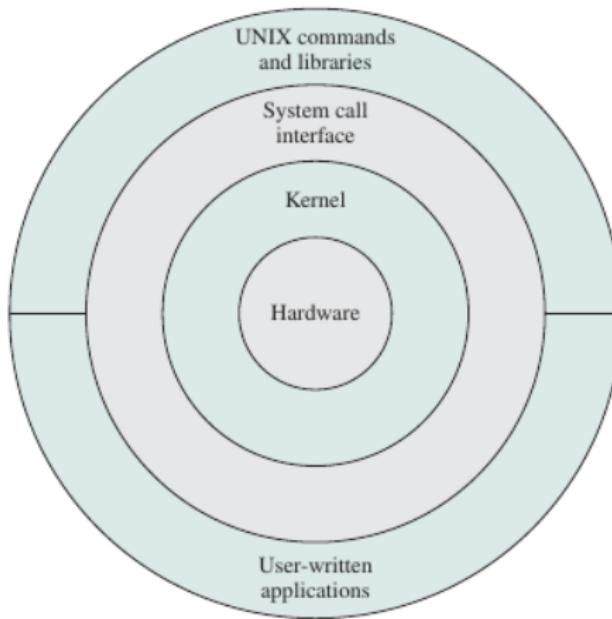
# Overview I



# Overview II

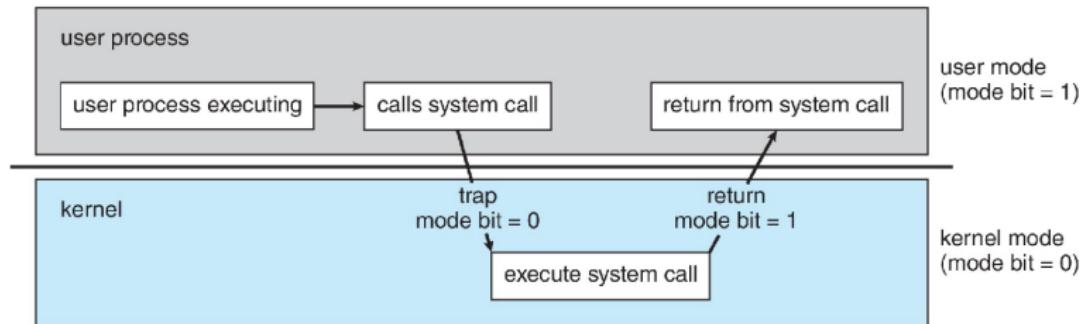


## Overview III



**Figure 2.16 General UNIX Architecture**

# Overview IV



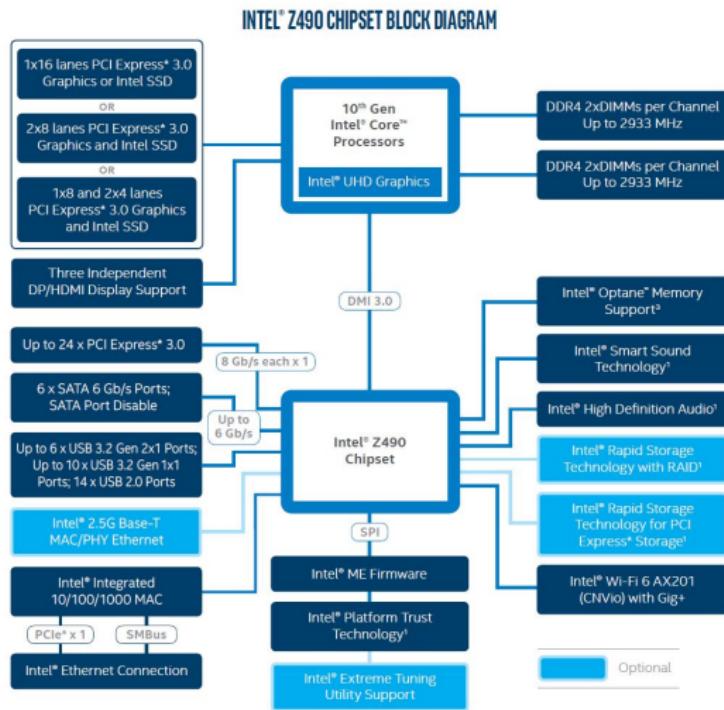
# Overview V

## EXAMPLES OF WINDOWS AND UNIX SYSTEM CALLS

The following illustrates various equivalent system calls for Windows and UNIX operating systems.

	Windows	Unix
Process control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File management	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device management	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communications	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shm_open() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()

# Hardware I



## Hardware II

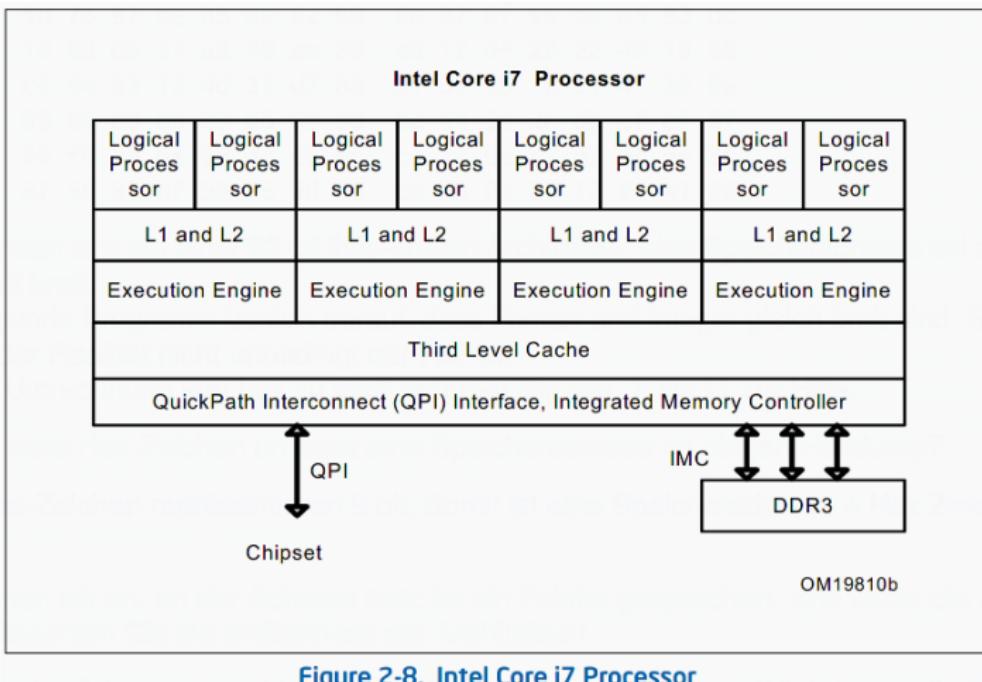
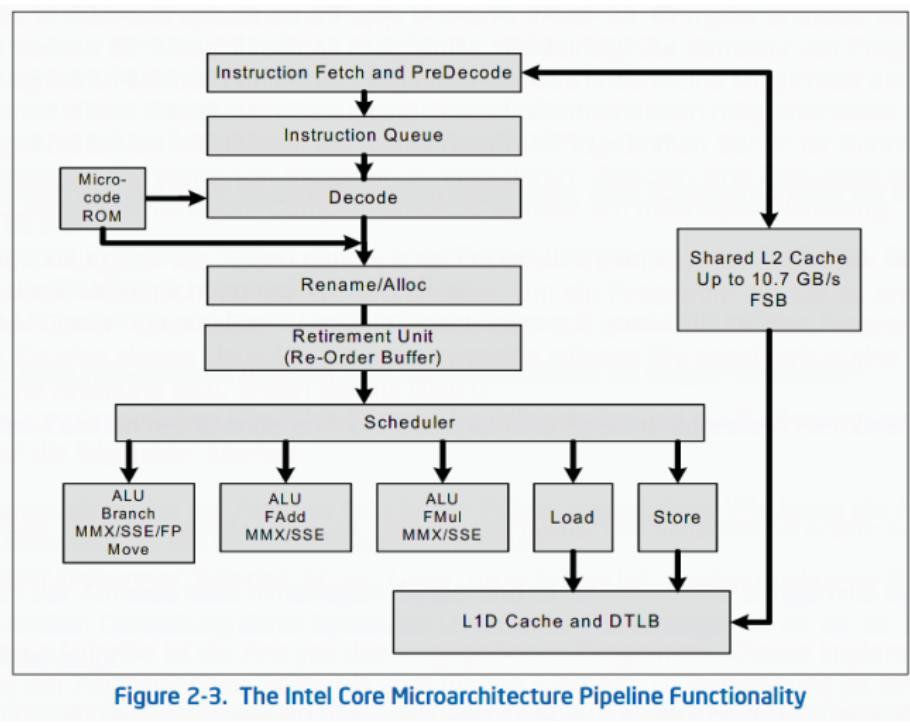


Figure 2-8. Intel Core i7 Processor

# Hardware III



# Hardware IV

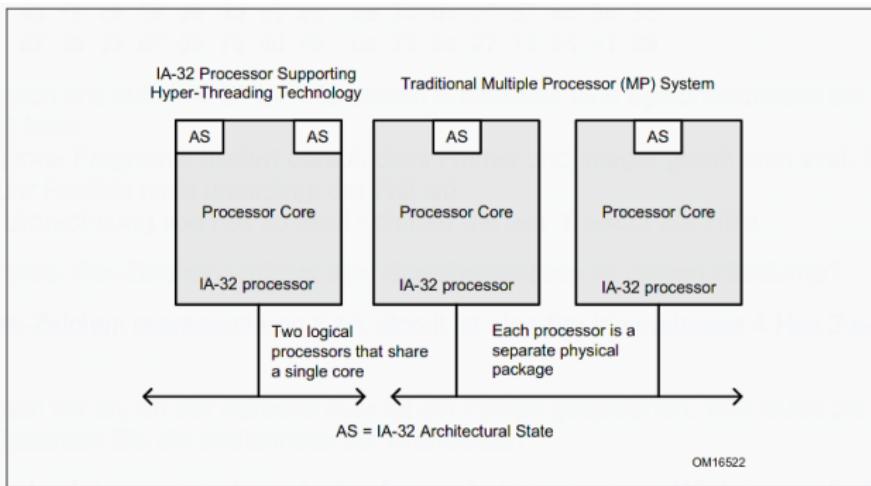
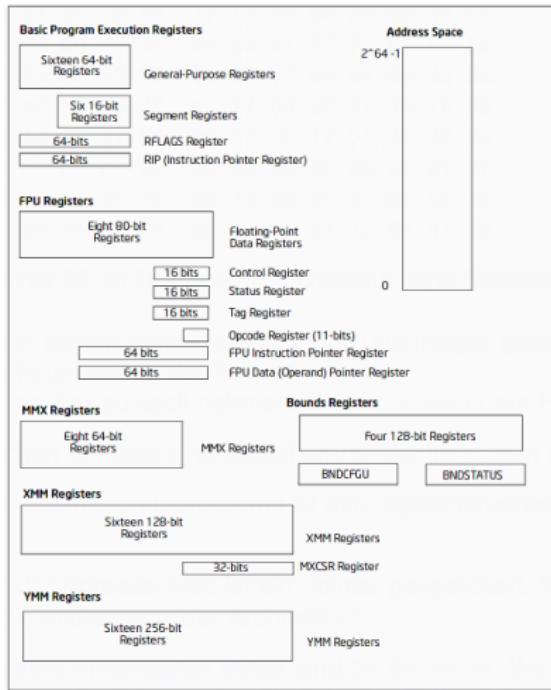


Figure 2-5. Comparison of an IA-32 Processor Supporting Hyper-Threading Technology and a Traditional Dual Processor System

# Hardware V



# Hardware VI

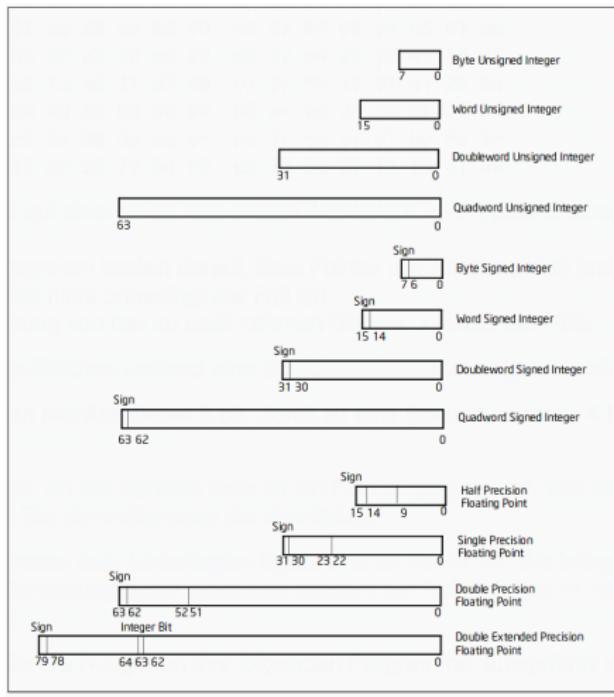
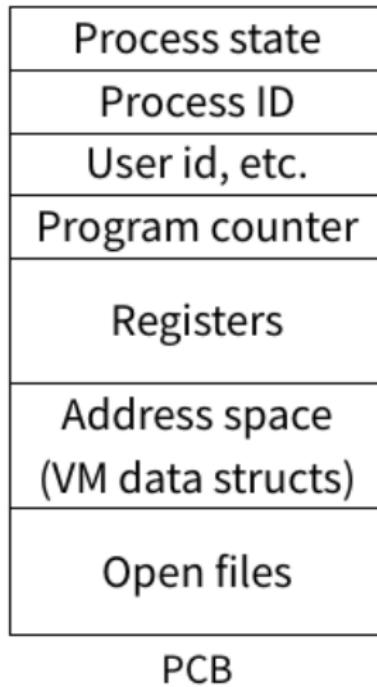
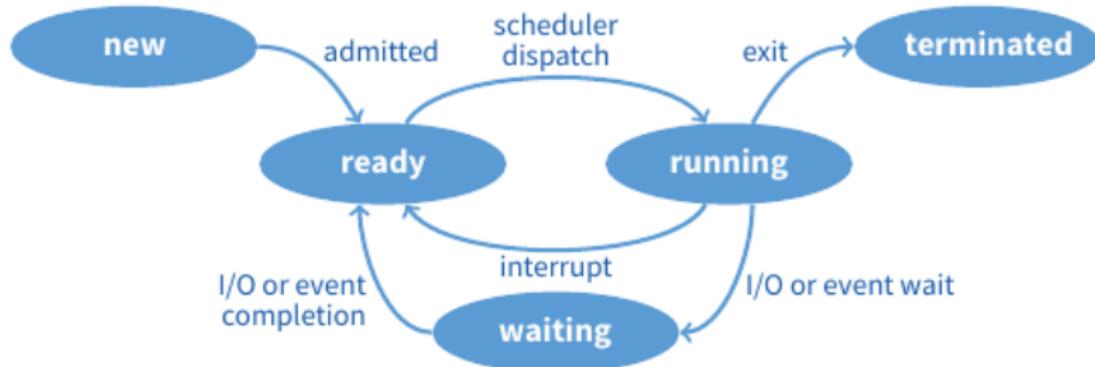


Figure 4-3. Numeric Data Types

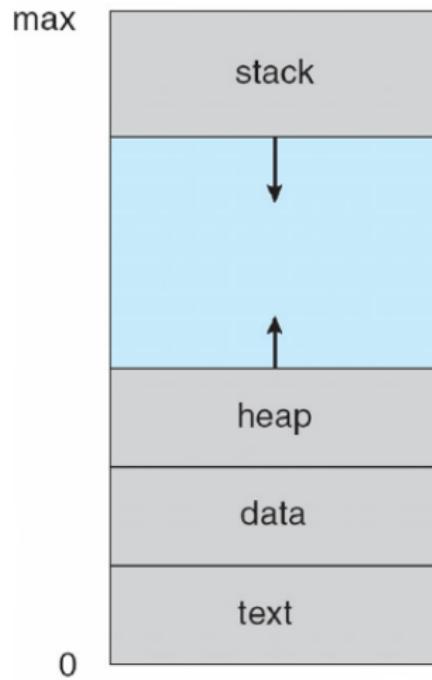
# Processes & Thread I



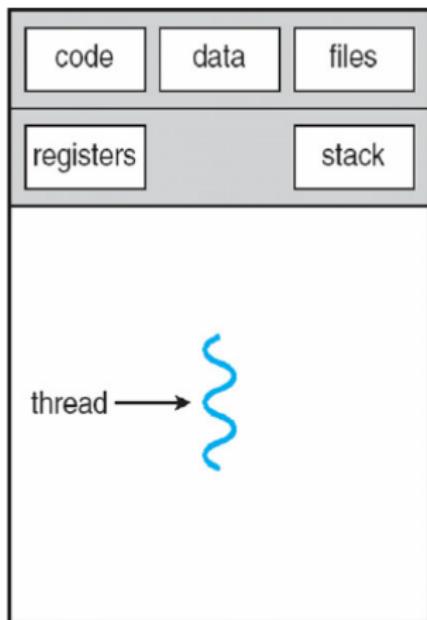
# Processes & Thread II



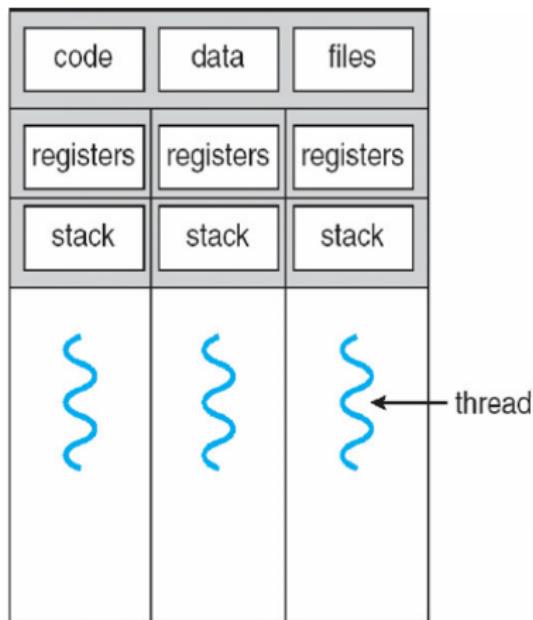
# Processes & Thread III



# Processes & Thread IV

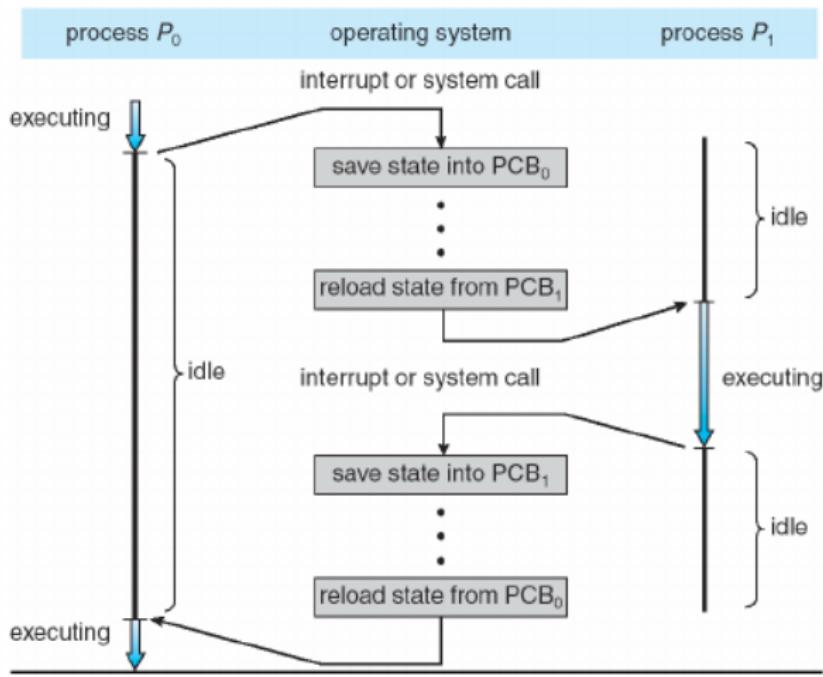


single-threaded process



multithreaded process

# Processes & Thread V



# Scheduling I

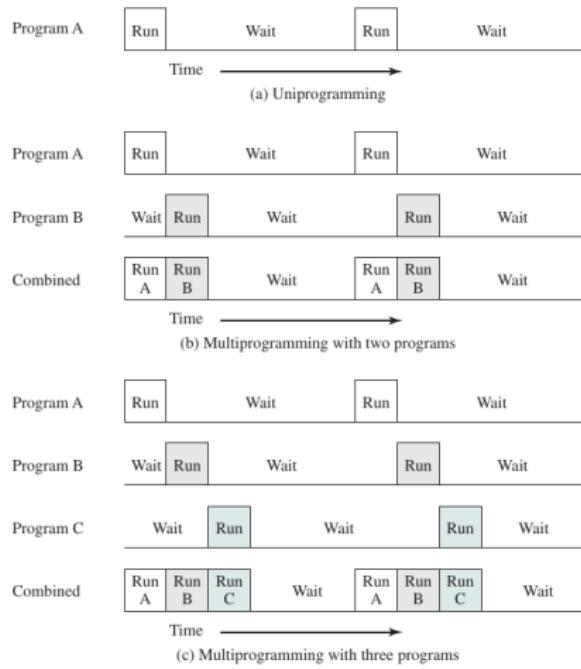


Figure 2.5 Multiprogramming Example

# Scheduling II

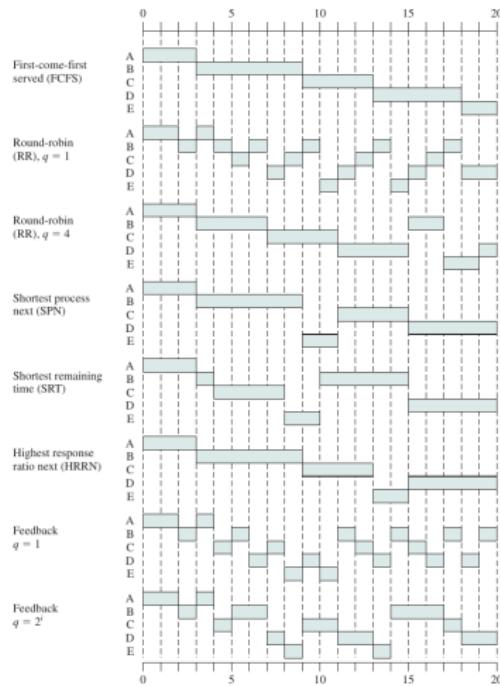
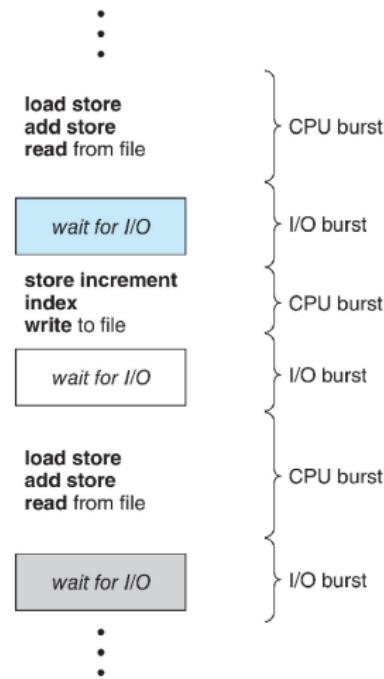
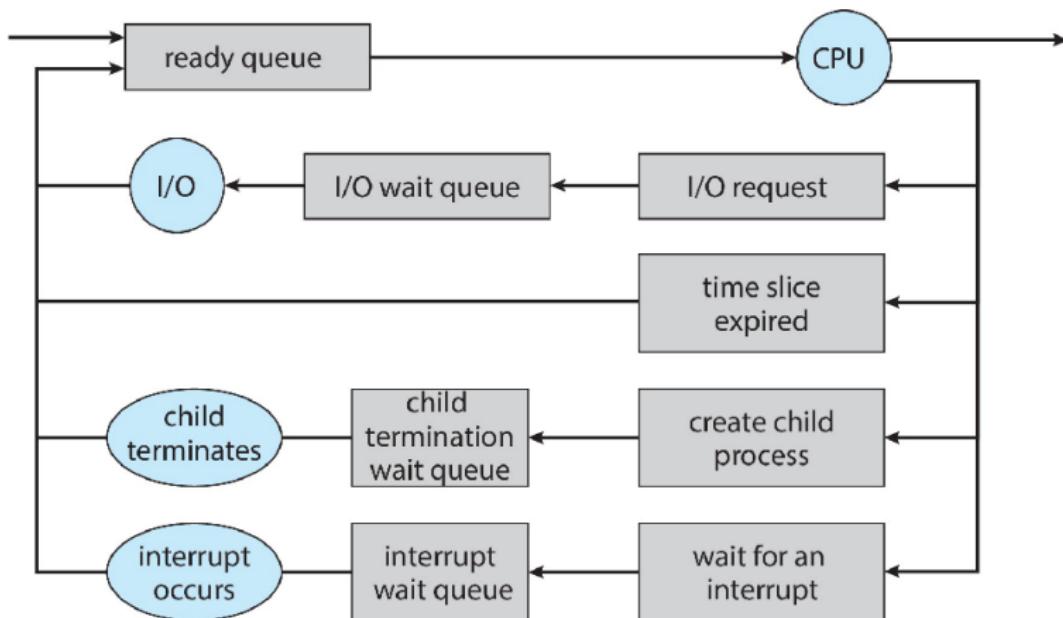


Figure 9.5 A Comparison of Scheduling Policies

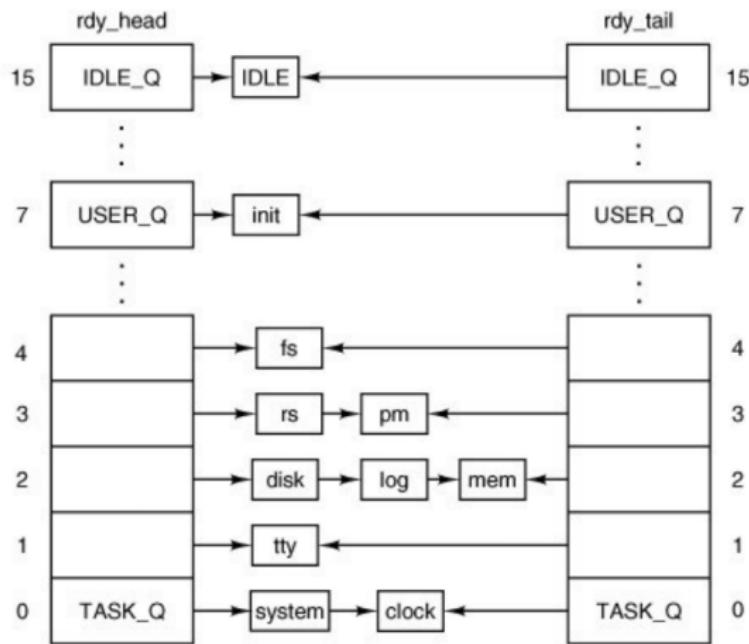
# Scheduling III



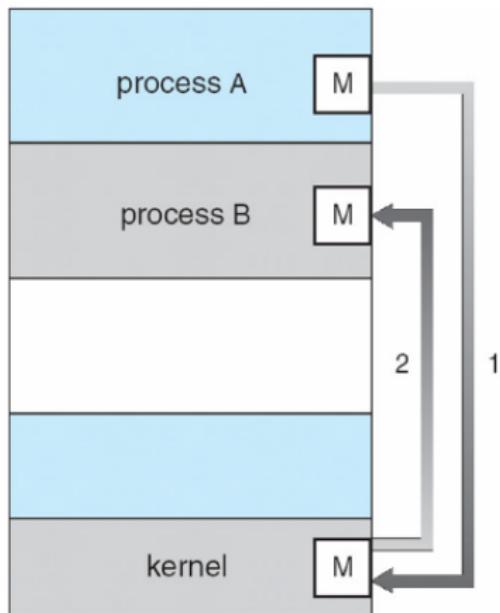
## Scheduling IV



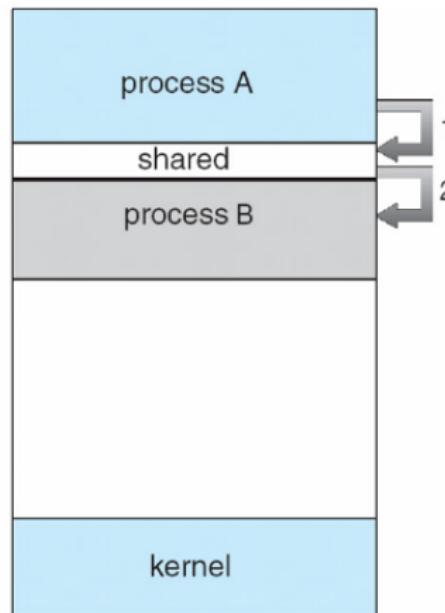
# Scheduling V



# Memory Management I

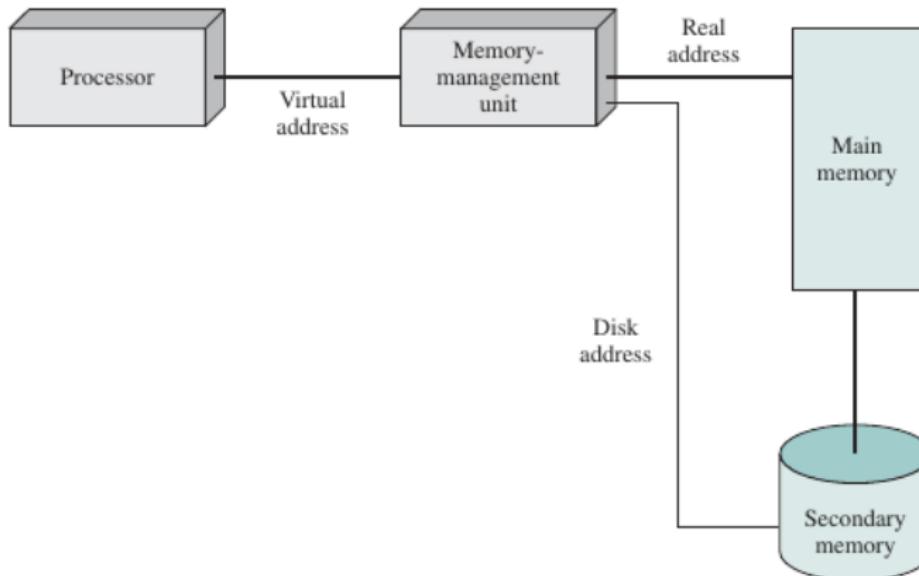


(a)



(b)

## Memory Management II



**Figure 2.10** Virtual Memory Addressing

# Memory Management III

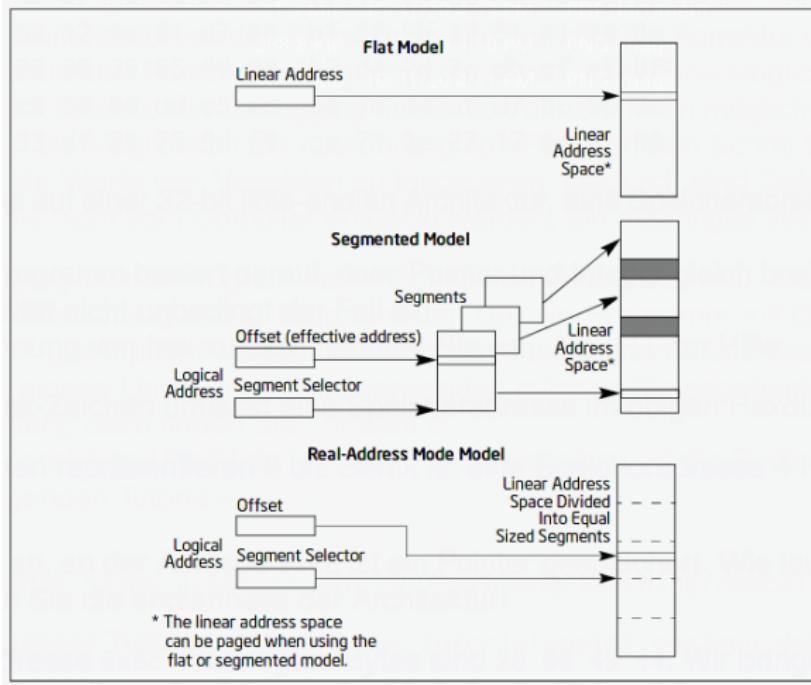
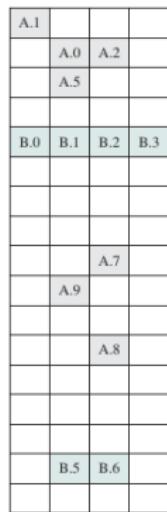


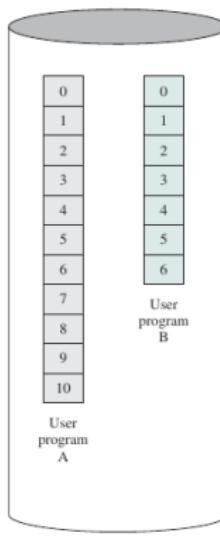
Figure 3-3. Three Memory Management Models

# Memory Management IV



Main memory

Main memory consists of a number of fixed-length frames, each equal to the size of a page. For a program to execute, some or all of its pages must be in main memory.

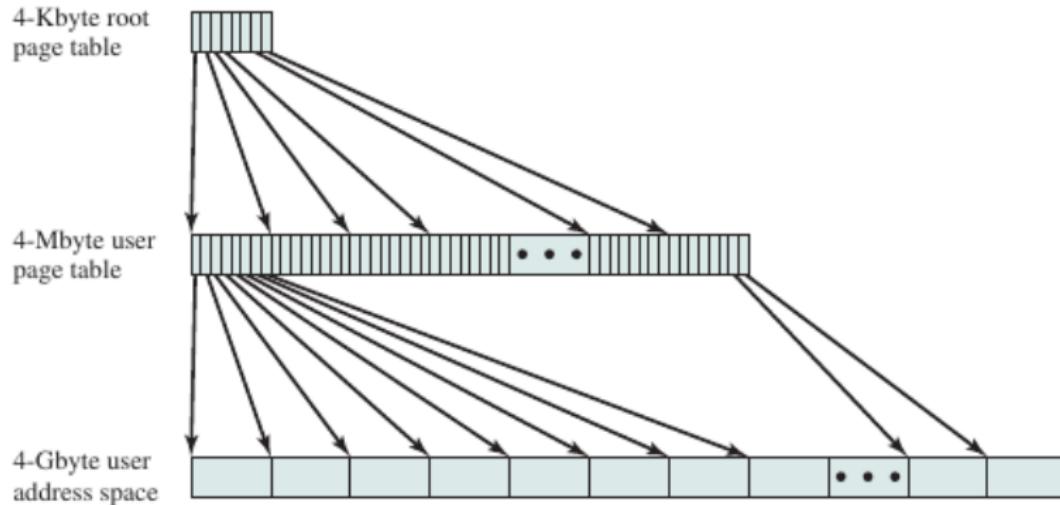


Disk

Secondary memory (disk) can hold many fixed-length pages. A user program consists of some number of pages. Pages for all programs plus the operating system are on disk, as are files.

**Figure 2.9** Virtual Memory Concepts

## Memory Management V



**Figure 8.4** A Two-Level Hierarchical Page Table

# Memory Management VI

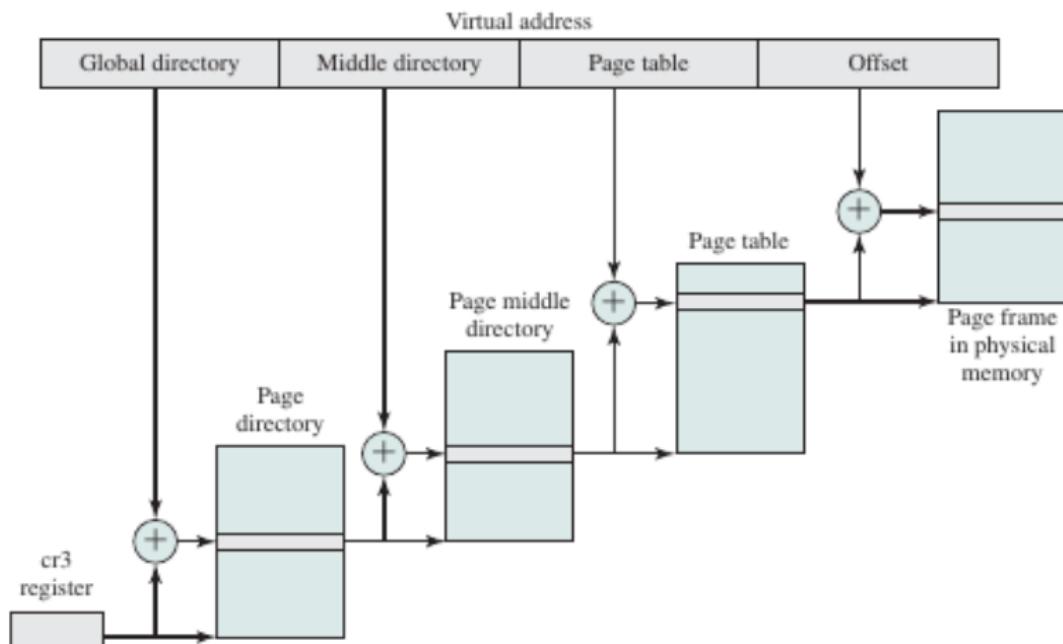


Figure 8.25 Address Translation in Linux Virtual Memory Scheme

# Memory Management VII

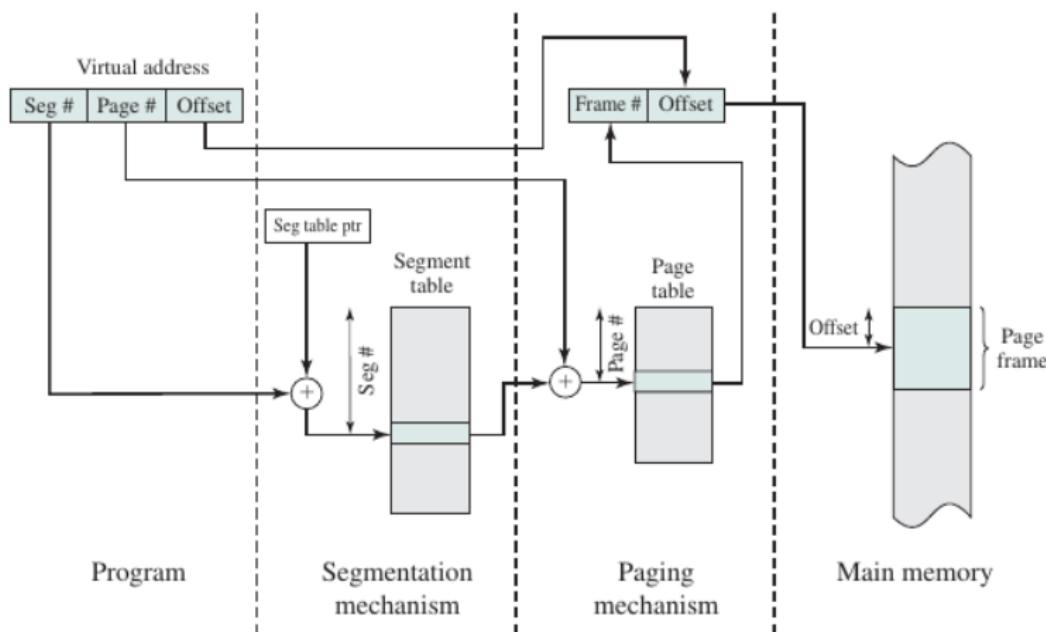
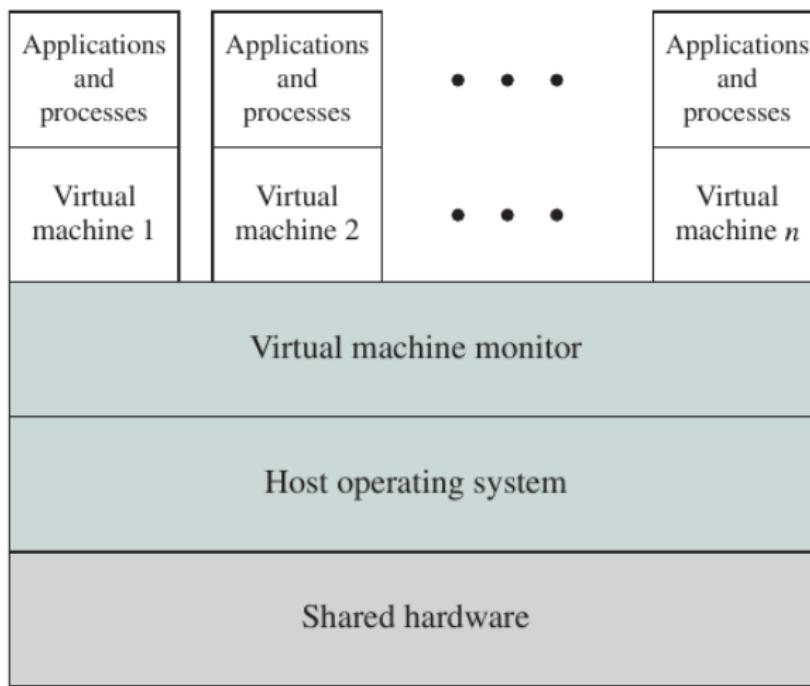


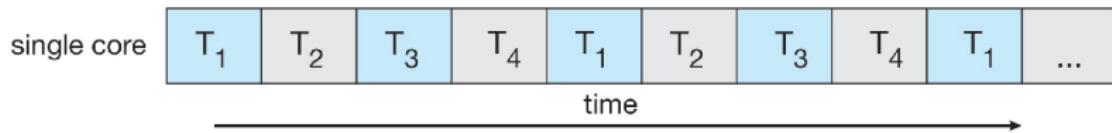
Figure 8.13 Address Translation in a Segmentation/Paging System

# Virtual Systems

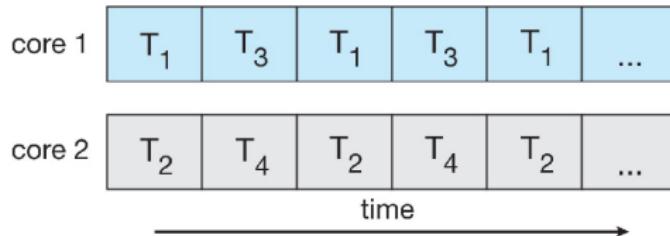


# Concurrency & Parallelism I

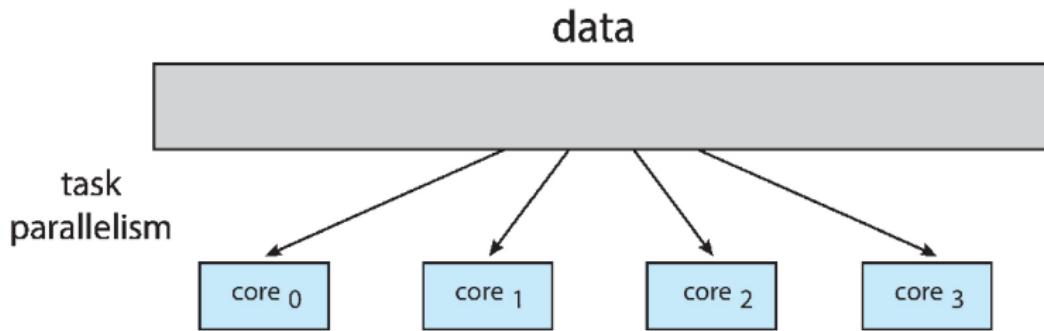
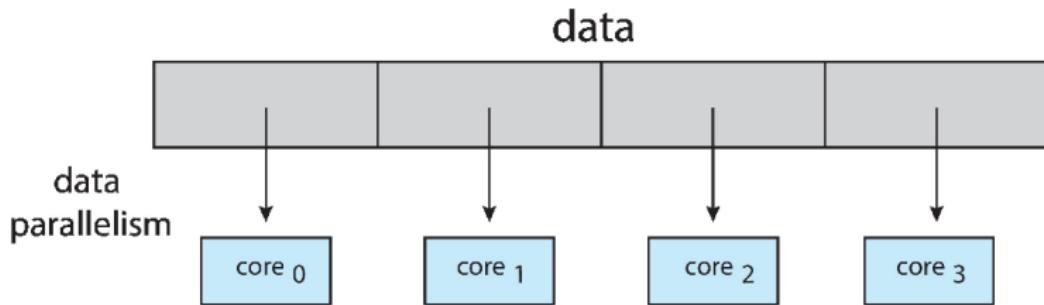
## § Concurrent execution on single-core system:



## § Parallelism on a multi-core system:



# Concurrency & Parallelism II



# Race Condition

## Process P1

- 
- 
- chin = getchar();
- 
- chout = chin;
- putchar(chout);
- 
- 

## Process P2

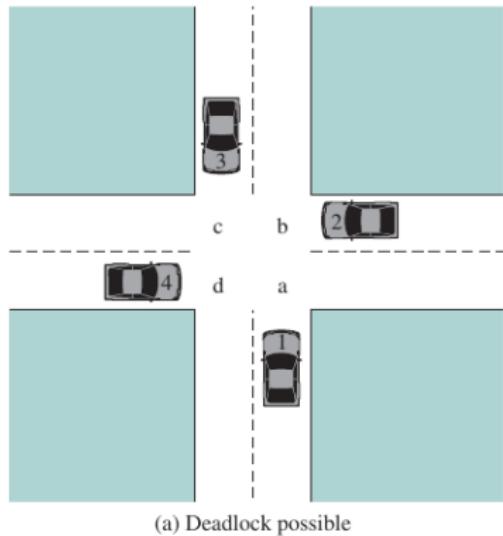
- 
- 
- chin = getchar();
- chout = chin;
- 
- putchar(chout);
-

# Synchronization

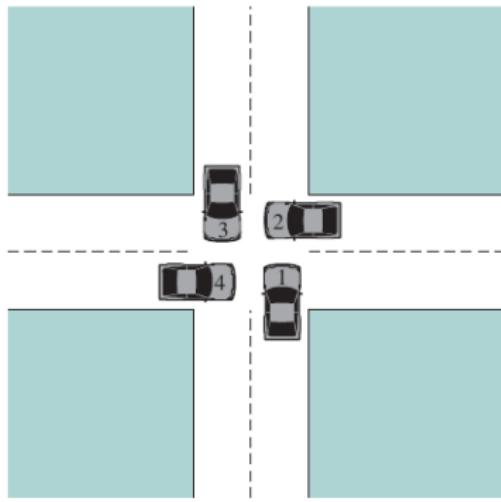
**Table 5.3** Common Concurrency Mechanisms

<b>Semaphore</b>	An integer value used for signaling among processes. Only three operations may be performed on a semaphore, all of which are atomic: initialize, decrement, and increment. The decrement operation may result in the blocking of a process, and the increment operation may result in the unblocking of a process. Also known as a <b>counting semaphore</b> or a <b>general semaphore</b> .
<b>Binary Semaphore</b>	A semaphore that takes on only the values 0 and 1.
<b>Mutex</b>	Similar to a binary semaphore. A key difference between the two is that the process that locks the mutex (sets the value to zero) must be the one to unlock it (sets the value to 1).
<b>Condition Variable</b>	A data type that is used to block a process or thread until a particular condition is true.
<b>Monitor</b>	A programming language construct that encapsulates variables, access procedures, and initialization code within an abstract data type. The monitor's variable may only be accessed via its access procedures and only one process may be actively accessing the monitor at any one time. The access procedures are <i>critical sections</i> . A monitor may have a queue of processes that are waiting to access it.
<b>Event Flags</b>	A memory word used as a synchronization mechanism. Application code may associate a different event with each bit in a flag. A thread can wait for either a single event or a combination of events by checking one or multiple bits in the corresponding flag. The thread is blocked until all of the required bits are set (AND) or until at least one of the bits is set (OR).
<b>Mailboxes/Messages</b>	A means for two processes to exchange information and that may be used for synchronization.
<b>Spinlocks</b>	Mutual exclusion mechanism in which a process executes in an infinite loop waiting for the value of a lock variable to indicate availability.

# Deadlocks I



(a) Deadlock possible



(b) Deadlock

**Figure 6.1 Illustration of Deadlock**

# Deadlocks II

Approach	Resource Allocation Policy	Different Schemes	Major Advantages	Major Disadvantages
Prevention	Conservative; undercommits resources	Requesting all resources at once	<ul style="list-style-type: none"> <li>Works well for processes that perform a single burst of activity</li> <li>No preemption necessary</li> </ul>	<ul style="list-style-type: none"> <li>Inefficient</li> <li>Delays process initiation</li> <li>Future resource requirements must be known by processes</li> </ul>
		Preemption	<ul style="list-style-type: none"> <li>Convenient when applied to resources whose state can be saved and restored easily</li> </ul>	<ul style="list-style-type: none"> <li>Preempts more often than necessary</li> </ul>
		Resource ordering	<ul style="list-style-type: none"> <li>Feasible to enforce via compile-time checks</li> <li>Needs no run-time computation since problem is solved in system design</li> </ul>	<ul style="list-style-type: none"> <li>Disallows incremental resource requests</li> </ul>
Avoidance	Midway between that of detection and prevention	Manipulate to find at least one safe path	<ul style="list-style-type: none"> <li>No preemption necessary</li> </ul>	<ul style="list-style-type: none"> <li>Future resource requirements must be known by OS</li> <li>Processes can be blocked for long periods</li> </ul>
Detection	Very liberal; requested resources are granted where possible	Invoke periodically to test for deadlock	<ul style="list-style-type: none"> <li>Never delays process initiation</li> <li>Facilitates online handling</li> </ul>	<ul style="list-style-type: none"> <li>Inherent preemption losses</li> </ul>

# Devices & Drivers I

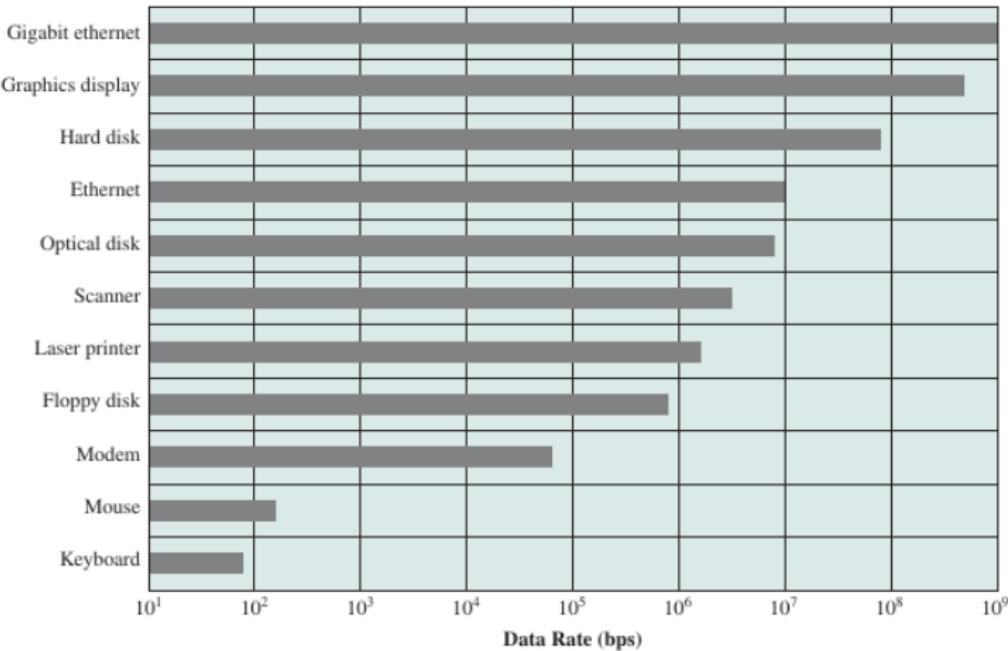
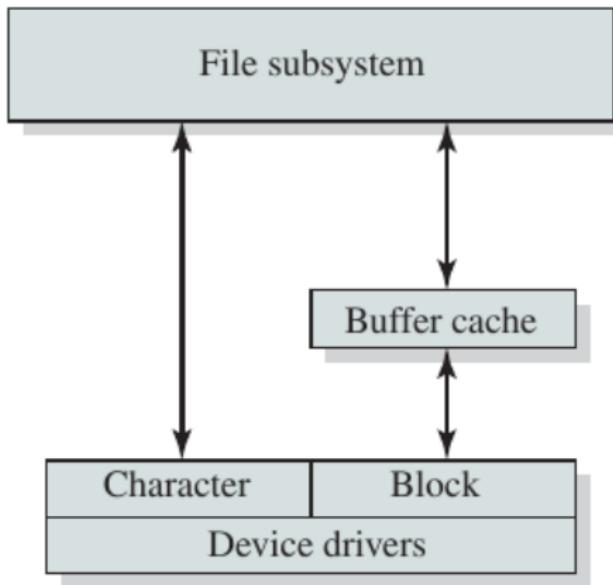


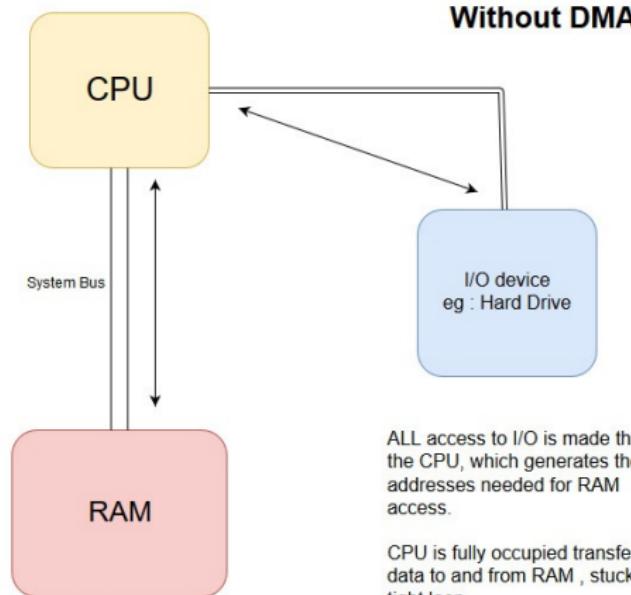
Figure 11.1 Typical I/O Device Data Rates

## Devices & Drivers II

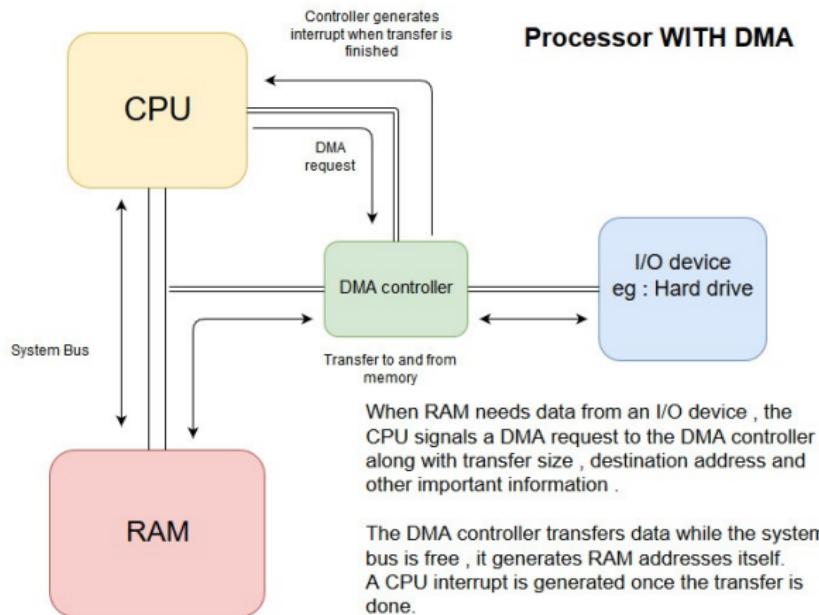


**Figure 11.12** UNIX I/O Structure

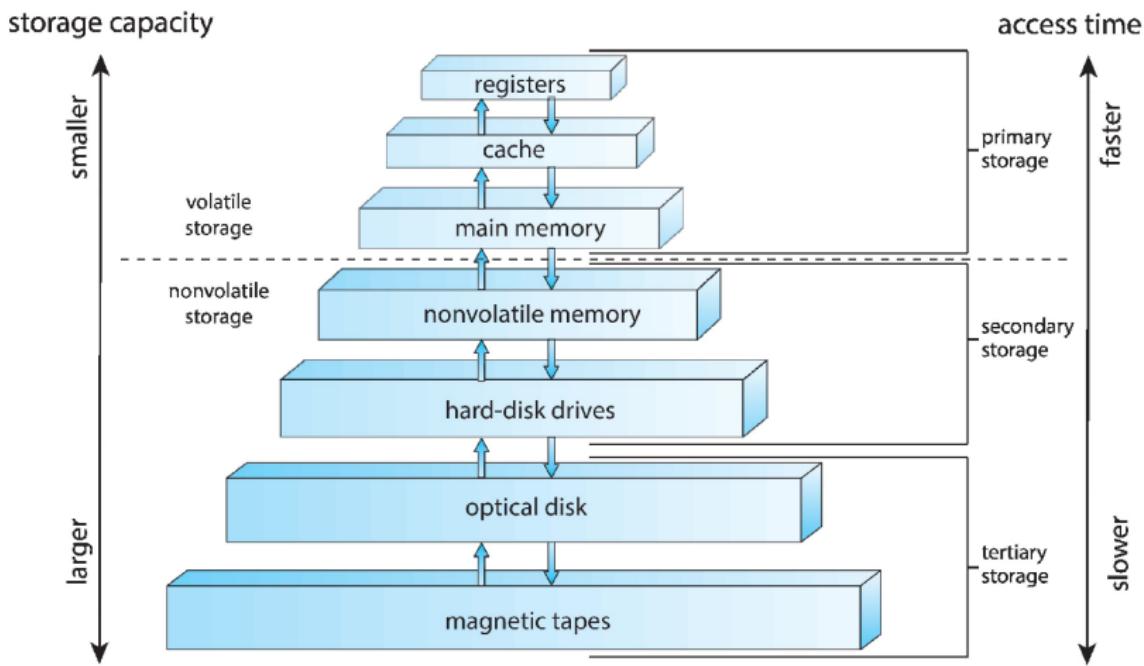
# Devices & Drivers III



# Devices & Drivers IV



# Devices & Drivers V

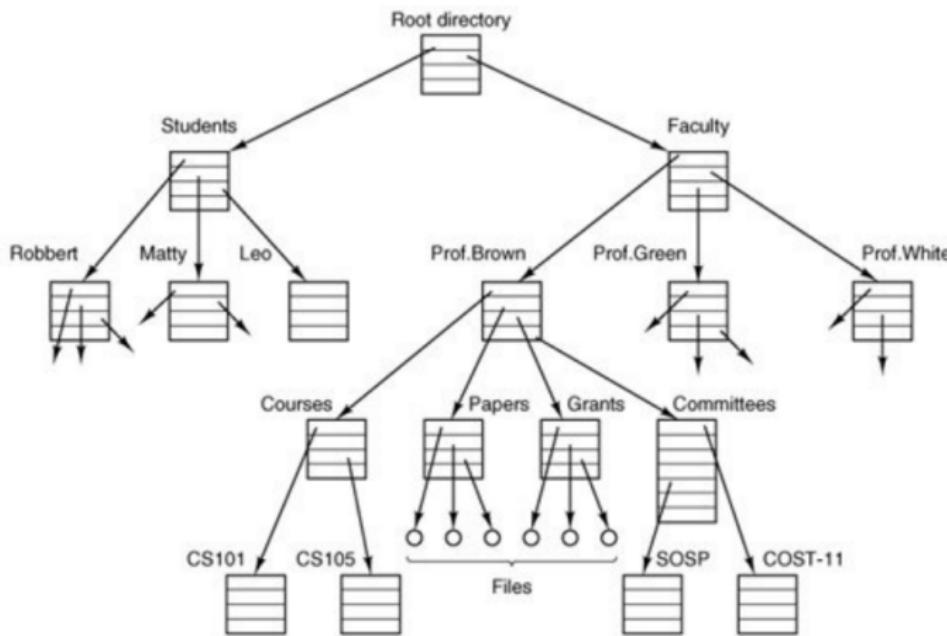


# Devices & Drivers VI

Level	1	2	3	4	5
Name	registers	cache	main memory	solid-state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25-0.5	0.5-25	80-250	25,000-50,000	5,000,000
Bandwidth (MB/sec)	20,000-100,000	5,000-10,000	1,000-5,000	500	20-150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape

Movement between levels of storage hierarchy can be explicit or implicit

# File Systems I



## File Systems II

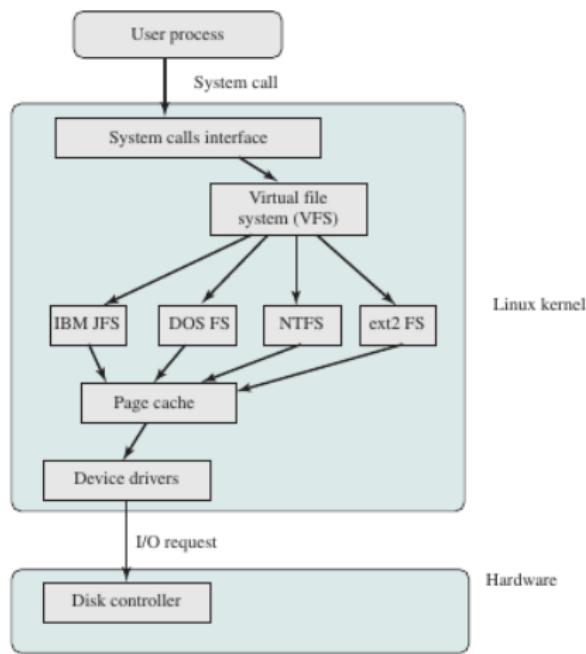
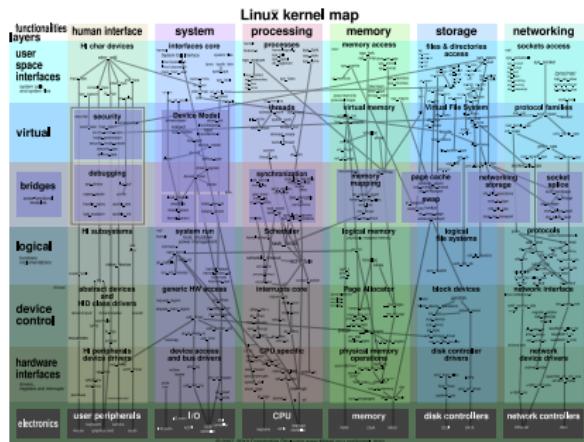


Figure 12.19 Linux Virtual File System Context

# All in One Map

Take 5 min to explore the map!



Take another 5 min to explore the source tree visually [on this site](#).

# C vs. Java

```
1 import java.util.Scanner;
2
3 public class Fakultaet
4 {
5     public static void main (String[] args)
6     {
7         Scanner scan = new Scanner(System.in);
8
9         int fakultaet = fak(scan.nextInt());
10        System.out.println("Fakultaet: " + fakultaet);
11    }
12    private static int fak(int x) {
13        return x <= 1 ? 1 : (x*fak(x-1));
14    }
15 }
```

```
1 #include <stdlib.h>
2 #include <stdio.h>
3
4 int fak(int);
5
6 int main(int argc, char *argv[]) {
7     char* buf = malloc(100 * sizeof(char));
8     fgets(buf, 100, stdin);
9
10    int fakultaet = fak(atoi(buf));
11    printf("Fakultaet: %d\n", fakultaet);
12    free(buf);
13 }
14
15 int fak(int x) {
16     return x <= 1 ? 1 : (x*fak(x-1));
17 }
```

# Contents of the Programming Part I

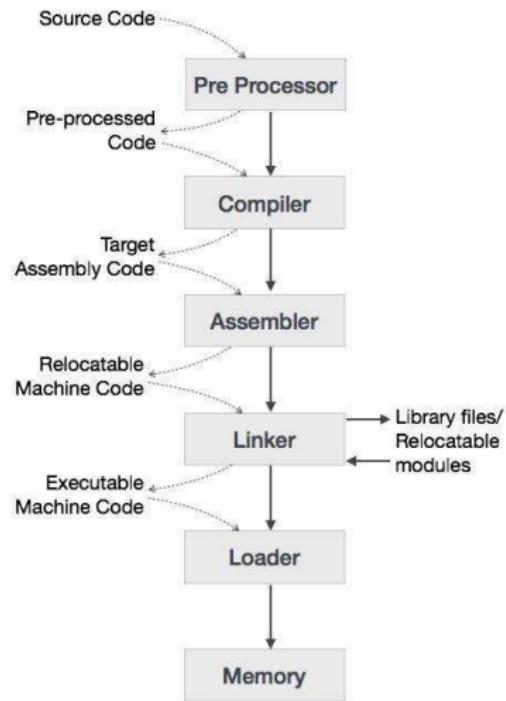
## 1. C Language

- 1.1 Basics
- 1.2 Pointers
- 1.3 Dynamic Memory Management
- 1.4 More on Types
- 1.5 Inside Malloc
- 1.6 Variables, Declarations & Scope
- 1.7 Large Program Organization

## 2. UNIX

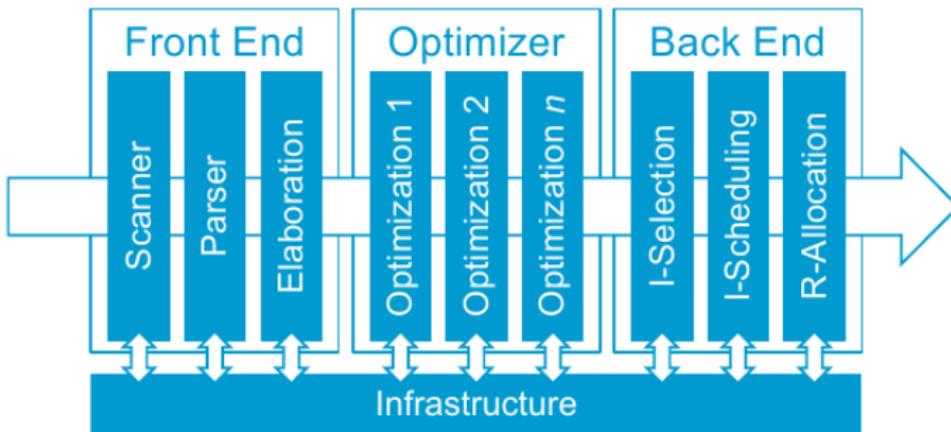
- 2.1 Arguments & Environment
- 2.2 Files & IO
- 2.3 Processes
- 2.4 Inter-Process Communication
- 2.5 Linking

# Compilers & Linkers I

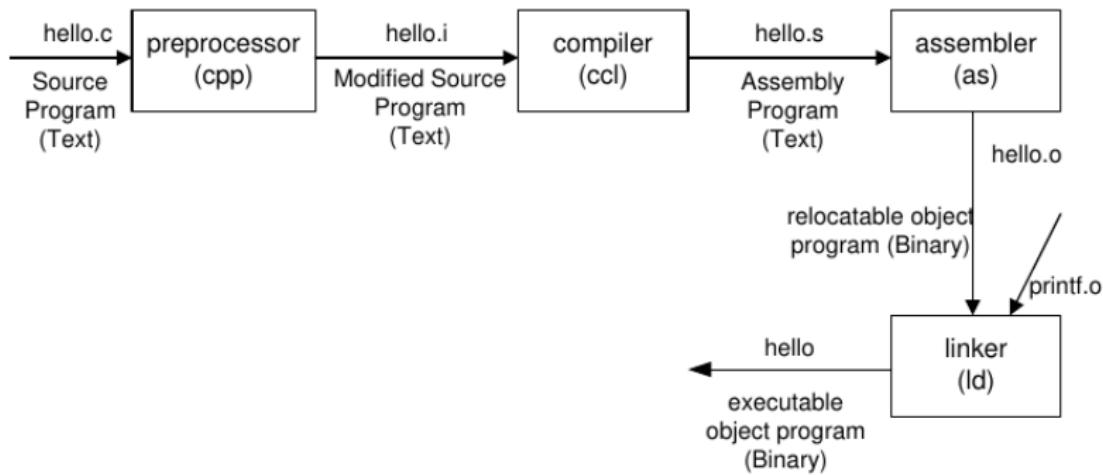


# Compilers & Linkers II

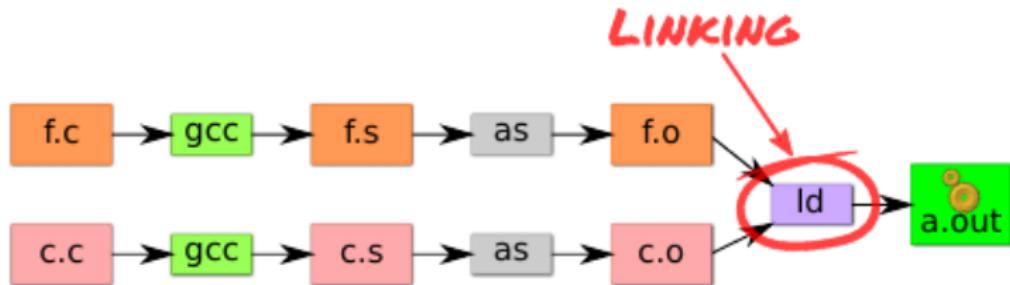
## Typical Structure of a Compiler



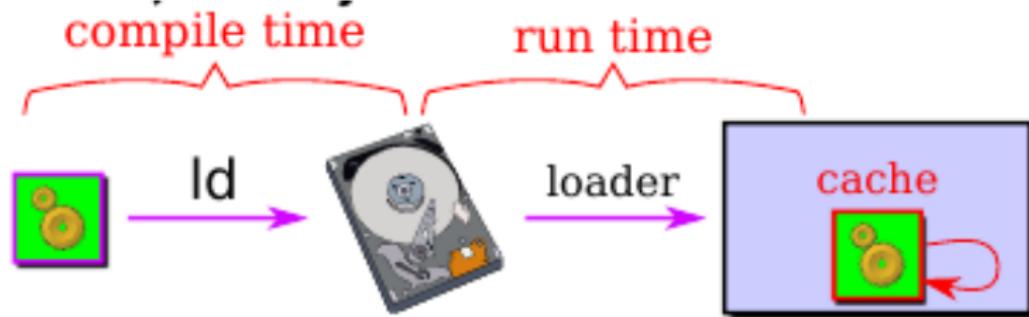
# Compilers & Linkers III



# Compilers & Linkers IV



## Compilers & Linkers V



# Compilers & Linkers VI

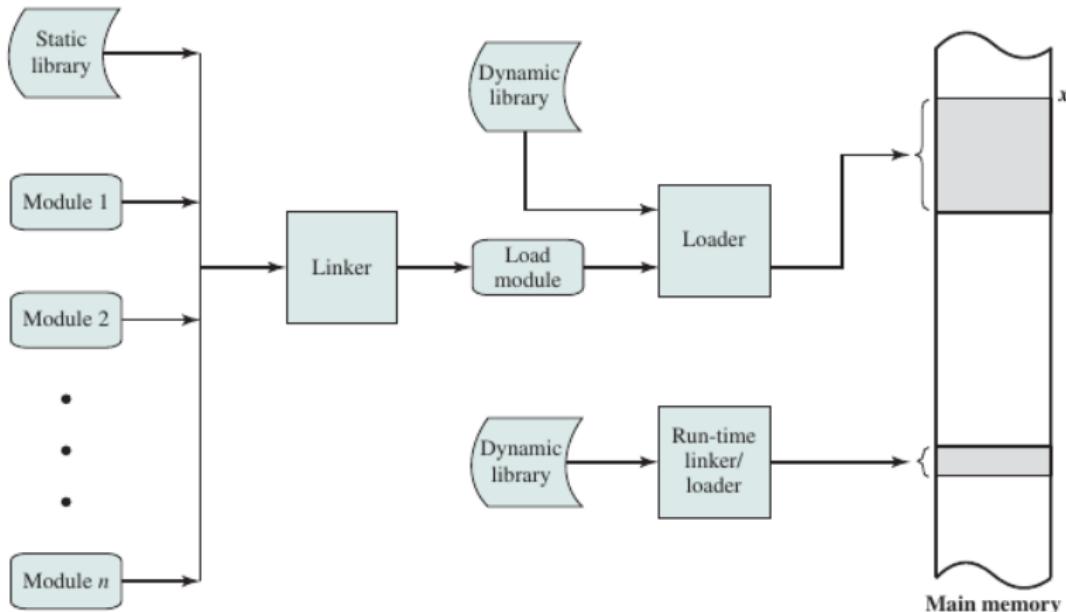


Figure 7.16 A Linking and Loading Scenario

# References |

-  W. Stallings, *Operating systems: internals and design principles*. Upper Saddle River, NJ: Pearson/Prentice Hall, 2009.
-  A. Silberschatz, P. B. Galvin, and G. Gagne, *Operating system principles*. John Wiley & Sons, 2006.
-  A. Rubini and J. Corbet, *Linux device drivers*. " O'Reilly Media, Inc.", 2001.
-  D. Mazieres. (Mar. 16, 2010), "Uni stanford cs 140: Operating systems," [Online]. Available: <https://www.scs.stanford.edu/10wi-cs140/> (visited on 11/02/2020).
-  A. S. Tanenbaum and H. Bos, *Modern operating systems*. Pearson, 2015.
-  Intel, "Intel 64 and ia-32 architectures software developers manual," , vol. 1,2,3,4, 2020.
-  L. Boulesteix. (Nov. 2, 2020), "What is the function of dma in a computer? - quora," [Online]. Available: <https://www.quora.com/What-is-the-function-of-DMA-in-a-computer> (visited on 11/02/2020).
-  C. Shulyupin. (Jul. 15, 2020), "Interactive map of linux kernel," [Online]. Available: <https://makelinux.github.io/kernel/map/> (visited on 11/02/2020).