

Operating Systems

Tutorial 2

Fabian Klopfer

16. November 2020

Intro

- ▶ Was the exercise sheet okay?
- ▶ How long did you take?
- ▶ Please tell me if something is odd with the screen/audio!
- ▶ Something unclear?

Section 1


Exercise Sheet 1

Exercise 1 I

1. What is the difference between kernel mode and user mode?

Kernel mode has full access to hardware, user mode is restricted (memory, hardware).

1. Kernel Mode: the machine operates with critical data structure, direct hardware (IN/OUT or memory mapped), direct memory, IRQ, DMA, and so on.
2. User Mode: users can run applications.



Kernel Mode "prevents" User Mode applications from damaging the system or its features.

Modern microprocessors implement in hardware at least 2 different states. For example under Intel, 4 states determine the PL (Privilege Level). It is possible to use 0,1,2,3 states, with 0 used in Kernel Mode.

Abbildung: The Linux Kernel Documentation Project: Fundamentals [1]

Exercise 1 II

2. Which of the following instructions should be allowed only in kernel mode, and why?

- ▶ Disabling all interrupts
⇒ **Kernel**, disables interaction HW → Kernel
- ▶ Reading the system clock
⇒ **User**, does not interfere with Kernel or other user programs
- ▶ Setting the system clock
⇒ **Kernel**, does interfere with Kernel and other user programs that use the system clock
- ▶ Changing the memory allocation tables
⇒ **Kernel**, interferes with the memory consistency of Kernel and all other programs

Exercise 1 III

3. **Modern operating systems virtualize memory address space, that is, they use logical memory addresses instead of physical (hardware) memory addresses. Describe two advantages of this approach.**
- ▶ Can use more memory than physically available (swapping)
 - ▶ Possible to have only necessary parts of program in-memory
 - ▶ Memory isolation: Every program has a separate logical memory space and no access to other addresses
 - ▶ Program memory doesn't need to be physically continuous (ease allocation)

Exercise 1 IV

4. When a user program reads from a disk file, it becomes blocked until the device driver has finished performing the read operation. Suppose the user program was not blocked when writing to a file, that is, execution of the program would continue while the device driver still performs the write operation. What problem could arise? Can the device driver solve it?

Exercise 1 V

Problems:

- ▶ What if another thread reads the file before the write?
Does it get the correct version of the contents? Need a signal for the driver to schedule access
- ▶ What if the write failed?
- ▶ What if the same program alters the buffer to be written during the write process?

Exercise 2 I

What is characteristic for the use of the following types of operating systems?

Operating systems for

- ▶ personal computers
Responsive, interactive, one user at a time
- ▶ embedded systems
Few resources, fixed set of programs, reliability, control over device
- ▶ sensor nodes
Few resources, often only one program to transfer/broadcast data
- ▶ real-time applications
Hard or soft deadlines for routines, preemptive scheduling, latency of program is fixed within an interval

Exercise 2 II

What is the difference between timesharing systems and multiprogramming systems?

- ▶ Timesharing systems: several users run programs on one computing system at the same time.
- ▶ Multiprogramming systems: users run several programs simultaneously.

⇒ Timesharing \subset Multiprogramming

Exercise 3 I

1. **Complete the table.** C.f. next slide
2. **Which quantities are typically specified with binary prefixes?**
Data rates: Transmission speed, Read and Write Speeds.
In general: Digital information.
3. **Why are there no binary prefixes in the lower part of the table?**
The bit is the smallest possible chunk of information. There are no „half“ bits.
4. **How large is 1 TB in TiB?**
 $\frac{10^{12}}{2^{40}} \sim 0.9094947017729282$

Exercise 3 II

	SI		binär			
	Name	Symbol	Name	Symbol	$b = 2$	$b = 1024$
24	yotta	Y	yobi	Yi	80	8
21	zetta	Z	zebi	Zi	70	7
18	exa	E	exbi	Ei	60	6
15	peta	P	pebi	Pi	50	5
12	tera	T	tebi	Ti	40	4
9	giga	G	gibi	Gi	30	3
6	mega	M	mebi	Mi	20	2
3	kilo	k	kibi	Ki	10	1
-3	milli	m				
-6	micro	μ				
-9	nano	n				
	...					

Tabelle: Completed table of prefixes from ex. 3

Fetch Decode Execute I

From Patterson & Hennessy [2]:

- A Fetch current instruction from memory by querying the memory with the program counter/ instruction pointer (RIP/EIP), advance it by the length of an instruction (32/64 bit)
- B Decode fetched instruction into micro-code and fetch stored registers from memory if necessary. Compute possible branch targets
- C Execute the micro-code instruction.
One of: Memory reference - add base and offset, register-register ALU instr. - compute what is spec'ed by op with values from registers, register-immediate ALU instr. - instr. defined second operand like increment or determine whether conditional branch condition is true.
- D Memory access: Write result to register or main memory

Fetch Decode Execute II




Abbildung: Illustration of instruction pipelining [3]

Fetch Decode Execute III

For unpipelined execution engines we have, with n instructions, k steps (here $k = 3$ for fetch decode and execute) and τ the time one instruction takes:

$$n \cdot k \cdot \tau$$

For pipelined execution engines we get, with i the number of steps before execution (here $i = 2$ for fetch and decode):

$$n \cdot \tau - i$$

Exercise 4 I

1. A CPU has a three-level fetch-decode-execute pipeline. Each stage of the pipeline takes the same time of 1 ns to process.

- 1.1 **How many instructions can the CPU process per second?**

$$\frac{1 \text{ instr.}}{1 \text{ ns}} = \frac{1 \text{ instr.}}{1 \text{ s} \cdot 10^{-9}} = 10^9 \cdot 1 \frac{\text{instr.}}{\text{s}}$$

1 billion instructions per second.

- 1.2 **How many instructions can the CPU process if the pipelines has five levels?**

Pipelining let's the elements of the FDE cycle operate asynchronous. The same amount, as in pipelines there is only one ALU/element that can execute the operations.

Exercise 4 II

2. A computer system has a cache (access time 2 ns), RAM (access time 10 ns) and an external storage device (access time 10 ms).

2.1 If the cache hit rate is 95 % and the main memory hit rate is 99 %, how long is the average access time?

$$0.95 \cdot 2 \text{ ns} + 0.05 \cdot (0.99 \cdot 10 \text{ ns} + 0.01 \cdot 10 \text{ ms}) = 5002.39 \text{ ns} \sim 5 \mu\text{s}$$

2.2 Which of the three access times dominates the average access time?

External memory as it is slower by orders of magnitude.

2.3 Another external memory device is connected. The measured average access time is now 2002.395 ns. What is the access time of the new external storage device?

$$2002.395 \text{ ns} = 0.95 \cdot 2 \text{ ns} + 0.05 \cdot (0.99 \cdot 10 \text{ ns} + 0.01 \cdot x \text{ ms})$$

Solving for x yields $x = 3\,996\,000 \text{ ns} \sim 4 \text{ ms}$

Exercise 5 I

1. Write a C program that outputs "Hello world!".

```
#include <stdio.h>

int main()
{
    printf("Hello world!\n");
    return 0;
}
```

Exercise 5 II

2. Compile your program using the flags

-std=c11 -g -Wall -Wpedantic. What does each of these flags mean?

- ▶ -std=c11 the 2011 ISO C standard
- ▶ -g debugging information
- ▶ -Wall turn on lots of error checking
- ▶ -Wpedantic strict ISO C conformity

Exercise 5 III

3. Which additional files are generated when you use the `-save-temps` flag and what do they contain?

- ▶ `.i` generated by the pre-processor, source with expanded pre-processor directives
- ▶ `.s` generated assembly code by the compiler
- ▶ `.o` generated machine code by the assembler
- ▶ `.out` generated executable by the linker

Exercise 6 I

Write a C program that outputs the sine function $\sin(x)$ for $x \in [0, 2\pi]$ on the console.

Exercise 6 II

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <math.h>
4
5  int main()
6  {
7      // screen resolution
8      const int WIDTH = 80;
9      const int HEIGHT = 25;
10
11     // scan resolution and rendering
12     const double WIDTH_STEP = 2.0 * M_PI / ((double) WIDTH);
13     const double HEIGHT_STEP = 2.0 / ((double) HEIGHT);
14     const double THICKNESS = 1.2;
15
16     // loop line by line, from left to right
17     for (double y = 1.0; y >= -1.0; y -= HEIGHT_STEP)
18     {
19         for (double x = 0.0; x <= 2.0 * M_PI; x += WIDTH_STEP)
20         {
21             double result = sin(x);
22             if (fabs(y - result) < THICKNESS * HEIGHT_STEP) {
23                 printf("*");
24             } else {
25                 printf(" ");
26             }
27         }
28         printf("\n"); // new line
29     }
30
31     // success
32     return 0;
33 }
```

Exercise 6 III

```
1  #include <stdio.h>
2  #include <math.h>
3  #include <string.h>
4
5  #define HEIGHT 16
6  #define WIDTH 64
7
8  int main () {
9      const float horizontal_step = 2 * M_PI / (WIDTH);
10     char grid[HEIGHT][WIDTH];
11
12     float x = 0;
13     int y;
14
15     memset(&grid, 0x20, sizeof(grid));
16
17     for (unsigned int i = 0; i < WIDTH; i++) {
18         y = (int) HEIGHT/2 * sin(x) + HEIGHT/2;
19         grid[y][i] = '*';
20         x += horizontal_step;
21     }
22
23     for (int i = HEIGHT - 1; i > -1; i--) {
24         printf("%.5s\n", (int) WIDTH, grid[i]);
25     }
26     return 0;
27 }
```

Section 2

Exercise sheet 2

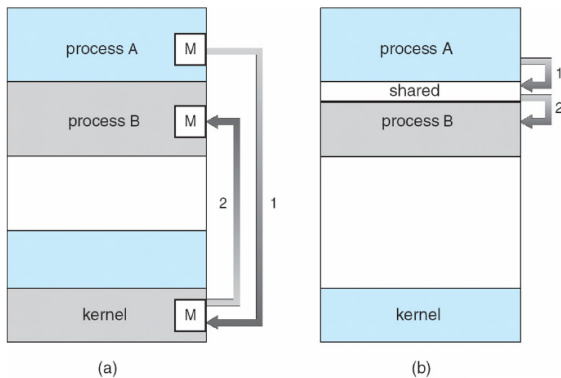


Abbildung: Visualization of how the kernel logically maps the main memory [4]

Process state
Process ID
User id, etc.
Program counter
Registers
Address space (VM data structs)
Open files

PCB

Abbildung: Elements contained in a process control block [5]




Abbildung: Elements contained in a process control block [4]

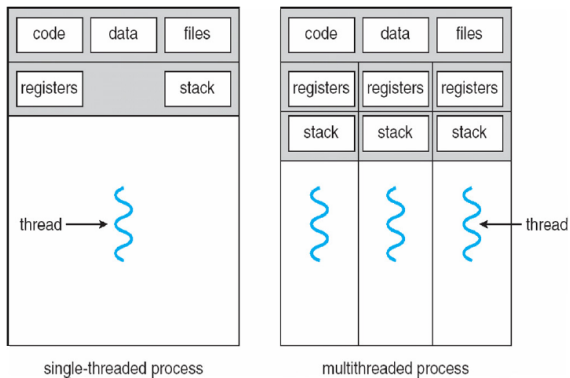


Abbildung: Elements contained in a process control block [4]

Section 3

C

Memory is uninitialized!

- ▶ <https://wiki.sei.cmu.edu/confluence/display/c/EXP33-C.+Do+not+read+uninitialized+memory>
- ▶ C Arrays

References I



(27. März 2003). „KernelAnalysis-HOWTO: Fundamentals“, Adresse: <https://tldp.org/HOWTO/KernelAnalysis-HOWTO-3.html> (besucht am 15.11.2020).



J. L. Hennessy und D. A. Patterson, *Computer architecture: a quantitative approach*. Elsevier, 2011.



(1. Nov. 2020). „Pipeline (Prozessor) Wikipedia“, Adresse: [https://de.wikipedia.org/wiki/Pipeline_\(Prozessor\)](https://de.wikipedia.org/wiki/Pipeline_(Prozessor)) (besucht am 15.11.2020).



A. Silberschatz, P. B. Galvin und G. Gagne, *Operating system principles*. John Wiley & Sons, 2006.



A. S. Tanenbaum und H. Bos, *Modern operating systems*. Pearson, 2015.



W. Stallings, *Operating systems: internals and design principles*. Upper Saddle River, NJ: Pearson/Prentice Hall, 2009.



(15. Nov. 2020). „Microsoft's Software Is Malware - GNU Project - Free Software Foundation“, Adresse: <https://www.gnu.org/proprietary/malware-microsoft.en.html> (besucht am 15.11.2020).



(15. Nov. 2020). „Apple's Operating Systems are Malware - GNU Project - Free Software Foundation“, Adresse: <https://www.gnu.org/proprietary/malware-apple.html#content> (besucht am 15.11.2020).